Wright State University

CORE Scholar

Kno.e.sis Publications

The Ohio Center of Excellence in Knowledge-Enabled Computing (Kno.e.sis)

3-26-2003

Ontology Driven Information Systems in Action (Capturing and Applying Existing Knowledge to Semantic Applications)

Amit P. Sheth Wright State University - Main Campus, amit@sc.edu

Follow this and additional works at: https://corescholar.libraries.wright.edu/knoesis

Part of the Bioinformatics Commons, Communication Technology and New Media Commons, Databases and Information Systems Commons, OS and Networks Commons, and the Science and Technology Studies Commons

Repository Citation

Sheth, A. P. (2003). Ontology Driven Information Systems in Action (Capturing and Applying Existing Knowledge to Semantic Applications). .

https://corescholar.libraries.wright.edu/knoesis/36

This Presentation is brought to you for free and open access by the The Ohio Center of Excellence in Knowledge-Enabled Computing (Kno.e.sis) at CORE Scholar. It has been accepted for inclusion in Kno.e.sis Publications by an authorized administrator of CORE Scholar. For more information, please contact library-corescholar@wright.edu.



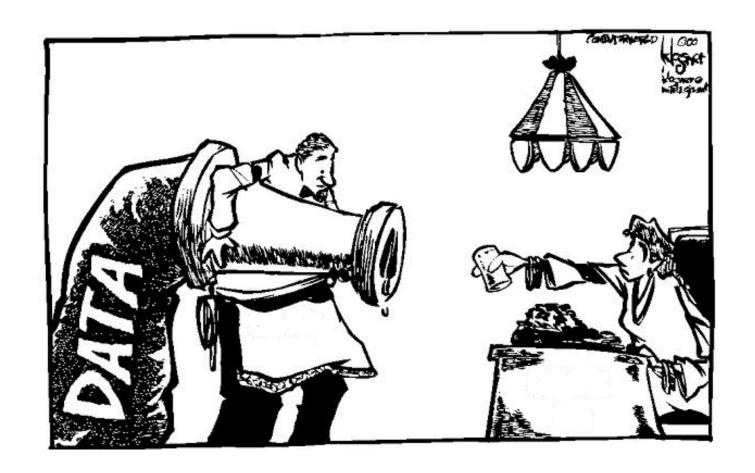


Invited Talk
"Sharing the Knowledge"
International CIDOC CRM Symposium
Washington DC, March 26 - 27, 2003

Amit Sheth

Semagix, Inc. and LSDIS Lab, University of Georgia

Syntax -> Semantics





Ontology-driven Information Systems are becoming reality

Software and practical tools to support key capabilities and requirements for such a system are now available:

- Ontology creation and maintenance
- Knowledge-based (and other techniques) supporting Automatic Classification
- Ontology-driven Semantic Metadata Extraction/Annotation and
 - Semantic normalization
- Utilizing semantic metadata and ontology
 - Semantic querying/browsing/analysis
 - Information and application integration

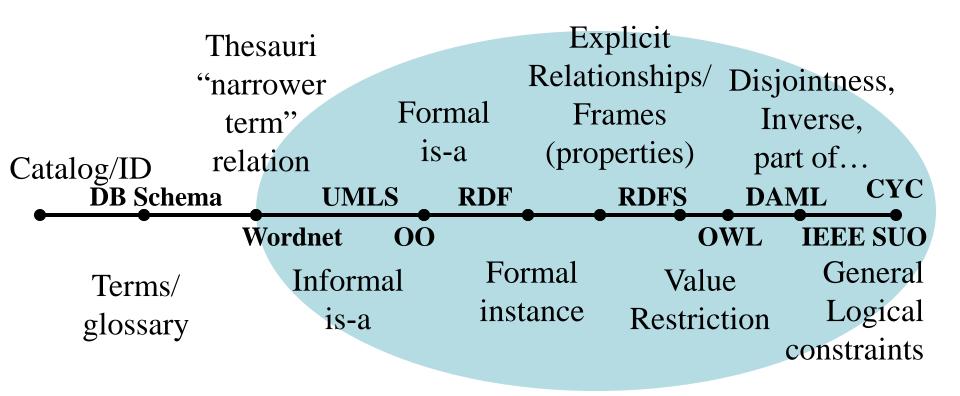
Achieved in the context of successful technology transfer from academic research (LSDIS lab, UGA's SCORE technology) into commercial product (Semagix's Freedom)

Ontology at the heart of the Semantic Web; Relationships at the heart of Semantics

Ontology provides underpinning for semantic techniques in information systems.

- A model/representation of the real world (relevant concepts, entities, attributes, relationships, domain vocabulary and factual knowledge, all connected via a semantic network). Basic of agreement, applying knowledge
- Enabler for improved information systems functionalities and the Semantic Web:
 - Relevant information by (semantic) Search, Browsing
 - Actionable information by (semantic) information correlation and analysis
 - Interoperability and Integration
- Relationships what makes ontologies richer (more semantic) than taxonomies ... see "<u>Relationships at the Heart of Semantic Web: Modeling,</u> <u>Discovering, Validating and Exploiting Complex Semantic Relationship</u>

Increasingly More Semantic Representation

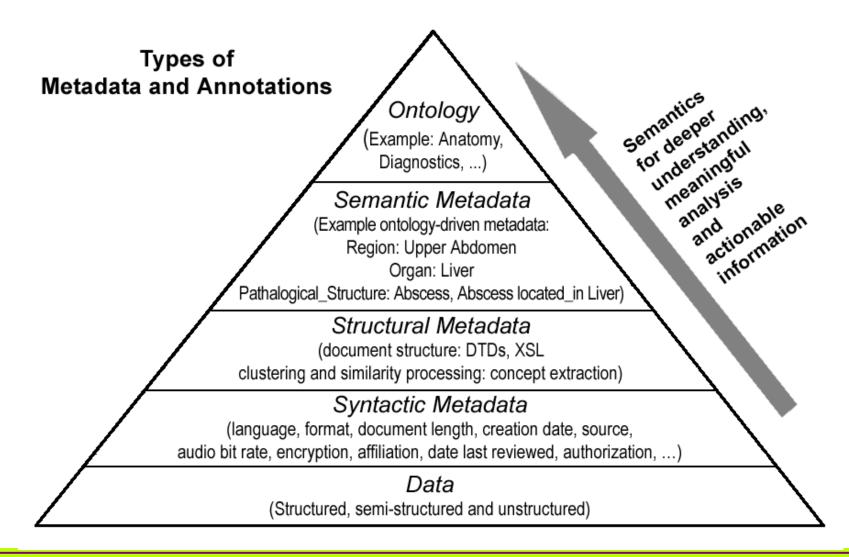


Simple Taxonomies

Expressive Ontologies
Better capability at higher complexity and computability



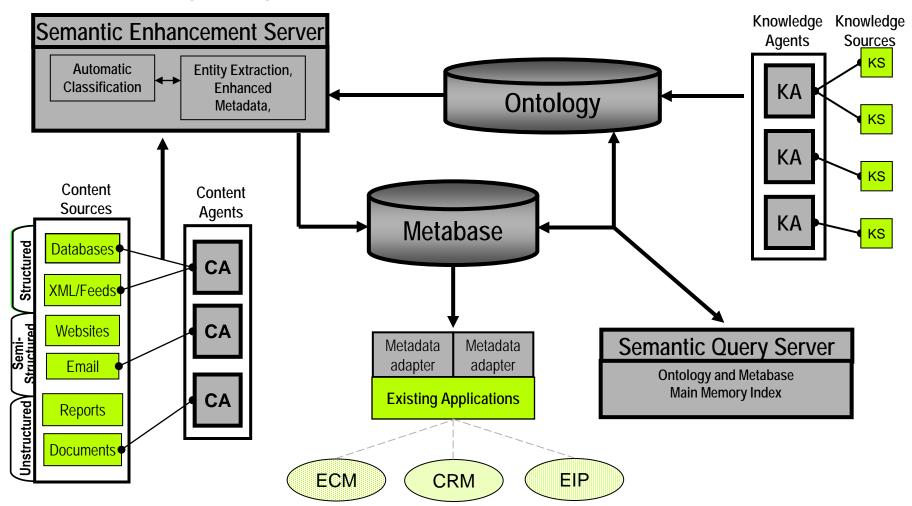
Metadata and Ontology: Primary Semantic Web enablers





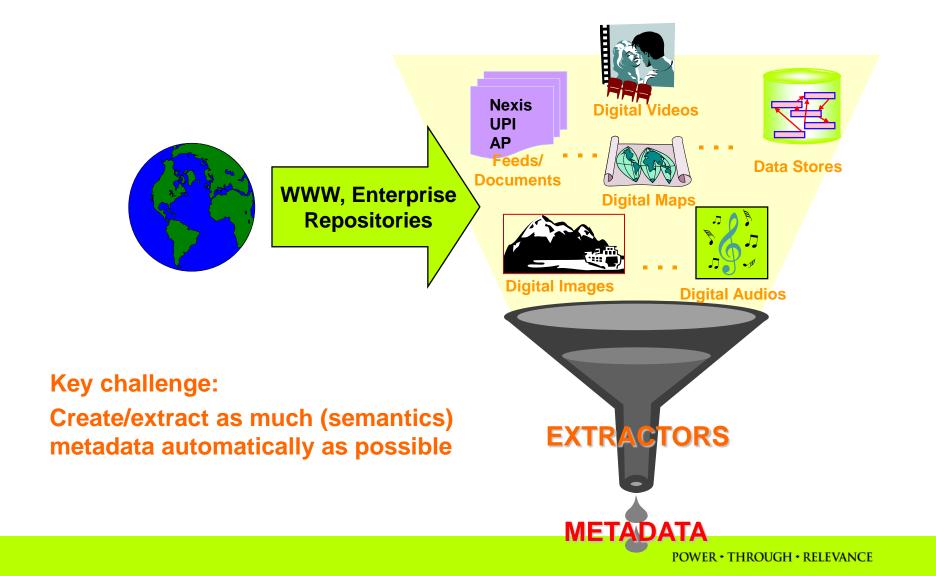
Semagix Freedom Architecture

(a platform for building ontology-driven information system)





Information Extraction and Metadata Creation



Automatic Classification & Metadata Extraction (Web page)



Braves refuse to offer Galarraga arbitration

Posted: Thursday December 07, 2000 6:15 PM

ATLANTA (AP) -- The Braves refused to offer salary arbitration to <u>Andres Galarraga</u> on Thursday, apparently ending the first baseman's career in Atlanta

Atlanta did offer arbitration to six of its former players who became free agents: pitchers <u>Andy Ashby</u>, <u>Terry Mulholland</u>, <u>John Burkett</u> and <u>Scott Kamieniecki</u>; first baseman <u>Wally Joyner</u> and

outfielder

Ashby ag year cont

free agent

Galarraga expired at Categorization

After missing the 1999 season because of cancer, Galar and 100 RBIs.

Auto

Free agents not offered arbitration by their former team until May 1.

The Braves made an offer Wednesday morning, but Ge said it was too low. Galarraga is seeking a two-year cor



Enter a URL:

http://sportsillustrateed.cnn.com/basebal/

Classify URL

Select a story from Virage:

galarraga



Classification Results

hhttp://sportsillustrated.cnn.com/baseball/ mlb/news/2000/12/07/galarraga_braves_ap/

Category	Predictors Agreement
baseball	80.36%
football	50.20%
golf	28.66%
business	21.91%
basketball	20.74%
hockey	20.54%
technology	19.55%
politics	12.01%
automotive	11.37%

Discovered Entities for Baseball Locations

Bonilla, Bobby Sportsperson Central (1266)

Joyner, Wally Sportsperson Atlanta (406)

Kamieniecki, Scott Sportsperson

Mulholland, Terry Sportsperson

Ashby, Andy Sportsperson

Galarraga, Andres Sportsperson

Semantic Metadata

Players offered arbitration have until Dec. 19 to accept or reject the oriers and can negotiate with their former teams through Jan. 8.

1

Ontology-directed Metadata Extraction (Semi-structured data)



Web Page



Enhanced Metadata Asset

Video Abstract	- Microsoft Internet Explorer	. 🗆 ×	
<u>File Edit Viev</u>	w Favorites Iools Help Address Links > Address Links > X	1	
	Next Update Update/Next Delete/Next Run Experts //rm.bbc.net.uk/news/olmedia/1320000/video/_1322301_milosevic01_anderson_vi.rm?title=	•	
Bitrat	te: 48 Format: real Width: 320 Height: 200 Invalidated: 0 Add		
Asset Id: Category: ExtractorName	31918736		
Keywords:			
Title:	Title: Kostunica pushes for war crimes law		
Surrogate:	http://news.bbc.co.uk/olmedia/1320000/images/_1324001_kostunica_afp300.jpg		
Media Type:	video		
URL:	pnm://rm.bbc.net.uk/news/olmedia/1320000/video/_1322301_milosevic01_anderson_vi.rm?title=		
Description:	The Yugoslav president says Belgrade will co-operate with the UN tribunal, which wants Slobodan Milosevic handed over.		
Clip Length:	230		
Parent URL:	http://news.bbc.co.uk/hi/english/world/europe/newsid_1324000/1324001.stm go there!		
Location:	: Belgrade, Yugoslavia, Europe		
People:	Yojislav Kostunica, Slobodan Milosevic		
Previous	Next Update Update/Next Delete/Next Run Experts	▼	
E	internet	- //	

Automatic Semantic Annotation of Text: Entity and Relationship Extraction



Blue-chip bonanza continues

Company Company Dow above 9,000 as HP, Home Depot lead advance; Microsoft upgrade helps techs. date time August 22, 2002: 11:44 AM EDT		
phrase By Alexandra Twin, CNN/Money Staff Writer		
New York (CNN/Money) - An upgrade of software leader Microsoft and strength in blue chips including company company Hewlett-Packard and Home Depot were among the factors pushing stocks higher at midday Thursday, with the Dow Jones industrial average spending time above the 9,000 level.		
Around 11:40 a.m. ET, the Dow Jones industrial average gained 65.06 to 9,022.09, continuing a more stock exchange than 1,300-point resurgence since July 23. The Nasdaq composite gained 9.12 to 1,418.37.		
The Standard & Poor's 500 index rose 9.61 to 958.97. company stockSym \$ \$ Hewlett-Packard (HPQ: up \$0.33 to \$15.03) Research, Estimates) said a report shows its share of		
the printer market grew in the second quarter, although another report showed that its share of the continent region continent computer server market declined in Europe, the Middle East and Africa. company stockSym \$ Home Depot (HD: up \$1.07 to \$33.75, Research, Estimates) was up for the third straight day after		
topping fiscal second-quarter earnings estimates on Tuesday.		
Tech stocks managed a turnaround. Software continued to rise after Salomon Smith Barney upgraded		
No. 1 software maker Microsoft (MSFT: up \$0.55 to \$52.83, Research, Estimates) to "outperform"		
from "neutral" and raised its price target to \$59 from \$56 Business software makers Oracle stockSym \$ \$ Company stockSym \$ \$ \$ \$ Company stockSym \$ \$ \$ \$ \$ \$ Company stockSym \$ \$ \$ \$ \$ \$ \$ \$ \$ \$ \$ \$ \$ \$ \$ \$ \$ \$ \$		
Research, Estimates) and BEA Systems (BEAS: up \$0.28 to \$7.12, Research, Estimates)		
all rose in tandem.		



Competitors

Automatic Semantic Annotation

<MediaType FormalName="Text" />

- < DataContent>

<MimeType FormalName="text/vnd.IPTC.NITF" />



<Property FormalName="Headquarters" Value="Houston, Texas, United States of American Ameri

<Property FormalName="Sector" Value="Energy" />

<Property FormalName="Industry" Value="Integrated Oil and Gas" />

Semantic Metadata Enhancement **Content Tags Semantic Metadata** Classification: Channel Partners, E-Business Solutions Company: Cisco Systems, Inc. **Content Tags Semantic Metadata Syntactic Metadata** Classification: Channel Partners, Producer: BusinessWire E-Business Solutions Source: Bloombera Date: Sept. 10 2001 Location: San Jose, CA URL: http://bloomberg.com/1.htm Media: Text E-Business Classification **Partners** Solutions **Semantic** Uniquely Metadata **Classification Committee** exploiting Extraction Knowledge-base, Machine Learning & (also syntactic) Statistical Techniques in the right Semantic Metadata **Content Tags** Company: Cisco Systems, Inc. Ontology Channel Partner E-Business Solution Classification: Channel Partners, E-Business Solutions Industry **Ticker** CIS-1270 Channel Partner: Siemens Network Security Channel Partner: Voyager Network Channel Partner: Siemens Network CIS-320 Channel Partner: Wipro Group Learning E-Business Solution: CIS-1270 Security CIS-6250 E-Business Solution: CIS-320 Learning Group **Finance** E-Business Solution: CIS-6250 Finance **Enabling powerful linking** E-Business Solution: CIS-1005 e-Market Ulvsvs CIS-1005 Group e-Market Ticker: CSCO of actionable information Industry: Telecommunication, . . . and facilitating important Sector: Computer Hardware Executives provider of Cisco semantic applications Executive: John Chambers Systems Competition: Nortel Networks such as knowledge Competition discovery and link competes with Syntactic Metadata analysis Producer: BusinessWire Source: Bloomberg Date: Sept. 10 2001 (user's task of manually Location: San Jose, CA retrieving all the information he URL: http://bloomberg.com/1.htm needs to know is greatly Media: Text minimized; he can spend more

Semantic Enhancement

minimized; he can spend more time making effective decisions)

XML content item with enriched semantic tagging,

ready to be queried

The CIDOC CRM can be an excellent starting point for building the Semantic Web and ontology-driven information system for exchange, interoperability, integration of data/information and knowledge in the area of scientific and cultural heritage.

Types of Ontologies (or things close to ontology)

- Upper ontologies: modeling of time, space, process, etc.
- Broad-based or general purpose ontology/nomenclatures: Cyc, CIRCA ontology (Applied Semantics), WordNet
- Domain-specific or Industry specific ontologies
 - News: politics, sports, business, entertainment
 - Financial Market
 - Terrorism
 - ♦ (GO (a nomenclature), UMLS inspired ontology, ...)
- Application Specific and Task specific ontologies
 - Anti-money laundering
 - Equity Research



Practical Questions (for developing typical industry and application ontologies)

- Is there a typical ontology?
 - Three broad approaches:
 - social process/manual: many years, committees
 - automatic taxonomy generation (statistical clustering/NLP): limitation/problems on quality, dependence on corpus, naming
 - Descriptional component (schema) designed by domain experts; Assertional component (extension) by automated processes
- How do you develop ontology (methodology)?
- People (expertise), time, money
- Ontology maintenance

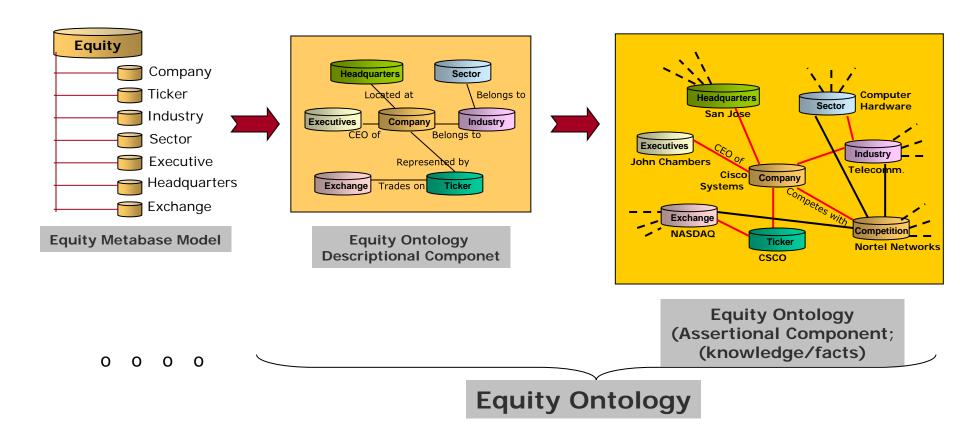


Practical Ontology Development Observation by Semagix

- Ontologies Semagix has designed:
 - Few classes to many tens (few hundreds) of classes and relationships (types); very small number of designers/knowledge experts; descriptional component (schema) designed with GUI
 - Hundreds of thousands to several millions entities and relationships (instances/assertions)
 - Tens of knowledge sources; populated by knowledge extractors
 - Primary scientific challenges faced: entity ambiguity resolution and data cleanup
 - Total effort: few person weeks



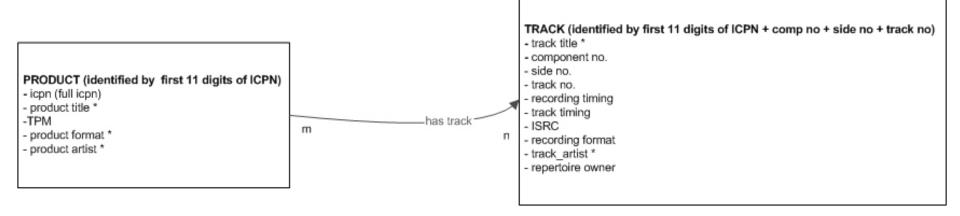
Ontology Example (Financial Equity domain)





Ontology with simple schema

- Ontology for a customer in Entertainment Industry
- Ontology Schema (Descriptional Component)
 - Only 2 high-level entity classes: Product and Track
 - A few attributes for each entity class
 - Only 1 relationship between the 2 classes: "has track"
 - Many-to-many relationship between the two entity classes
 - A product can have multiple tracks
 - A track can belong to multiple products





Entertainment Ontology Schema (Assertional Component)

- About 400K entity instances in ontology
- About 3.8M attribute instances in ontology
- Entity instances and attribute instances extracted by Knowledge Agents from 5 disparate databases
- Databases contain little overlapping and mostly 'dirty' data (unfilled values, inconsistent data)





Technical Challenges Faced

- Extremely 'dirty' data
 - Inconsistent field values
 - Unfilled field values
 - Field values appearing to mean the same, but are different
- Non-normalized Data
 - Same field value referred to, in several different ways
- Upper case vs. Lower case text analysis
- Modelling the ontology so that appropriate level (not too much, not too less) of information is modelled
- Optimizing the storage of the huge data
 - How to load it into Freedom (currently distributed across 3 servers)
- Scoring and pre-processing parameters changed frequently by customer, necessitating constant update of algorithm
- Efficiency measures



Effort Involved

Ontology Schema Build-Out (descriptional component)

Essentially an iterative approach to refining the ontology schema based on periodic customer feedback

 Very little technical effort (hours), but due to iterative decision making process with the multi-national customer, overall finalization of ontology took 3-4 weeks to complete

Ontology Population (assertional component/knowledge base)

- 5 Knowledge Agents, one for each database
- Automated ontology population using Knowledge Agents took <u>no longer</u>
 <u>than a day</u> for all the Agents

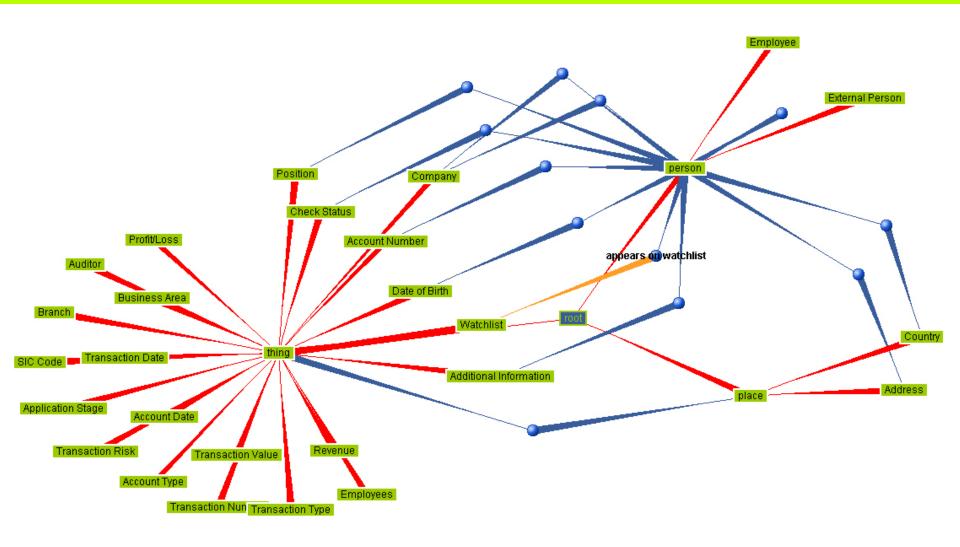


Example of Ontology with complex schema

- Ontology for Anti-money Laundering (AML) application in Financial Industry
- Ontology Schema (Descriptional Component)
 - About 40 entity classes
 - About 100 attribute types
 - About 50 relationship types between entity classes

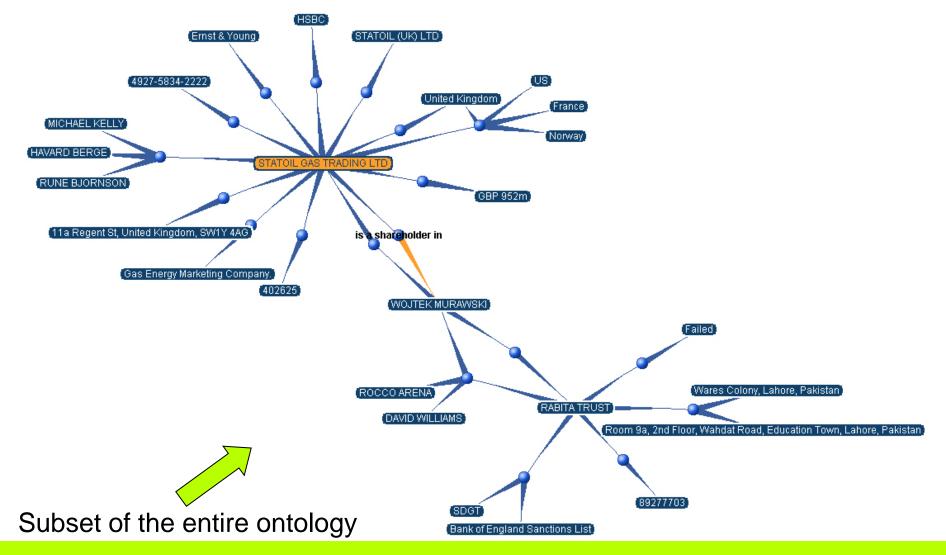


AML Ontology Schema (Descriptional Component)





AML Ontology Schema (Assertional Component)





AML (Anti-Money Laundering) Ontology

Ontology Schema (Assertional Component)

- About 1.5M entities, attributes and relationships
- 4 different sources for knowledge extraction
 - Dun and Bradstreet
 - ◆Corporate 192
 - Companies House
 - Hoovers

Effort Involved

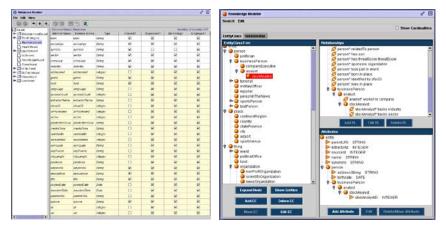
- Ontology schema design: 3 days
- Automated Ontology population using Knowledge Agents: 2 days



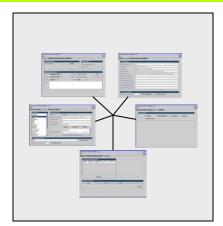
Technical Challenges Faced

- Complex ambiguity resolution at entity extraction time
- Modelling the ontology so that appropriate level (not too much, not too less) of information is modelled
- Knowledge extraction from sources that needed extended cookie/HTTPS handling
- Programming ontology modelling through API
- Chalking out a balanced risk algorithm based on numerous parameters involved

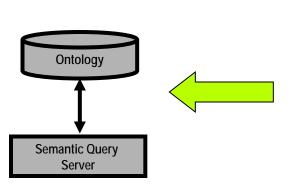
Ontology Creation and Maintenance Steps



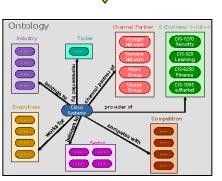
1. Ontology Model Creation



2. Knowledge Agent Creation



4. Querying the Ontology

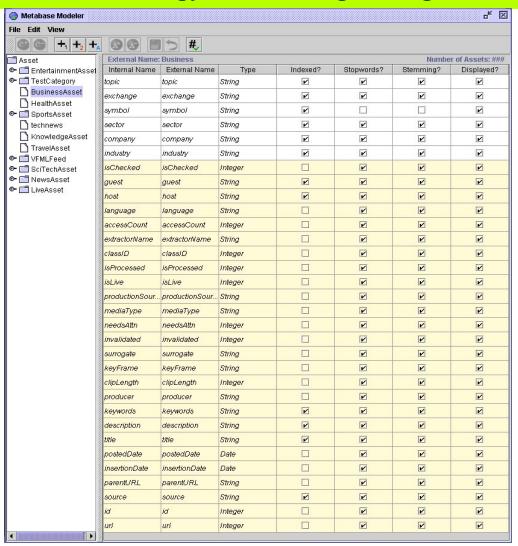


3. Automatic aggregation of Knowledge



Step 1: Ontology Model Creation

Create an Ontology Model using Semagix Freedom Toolkit GUIs

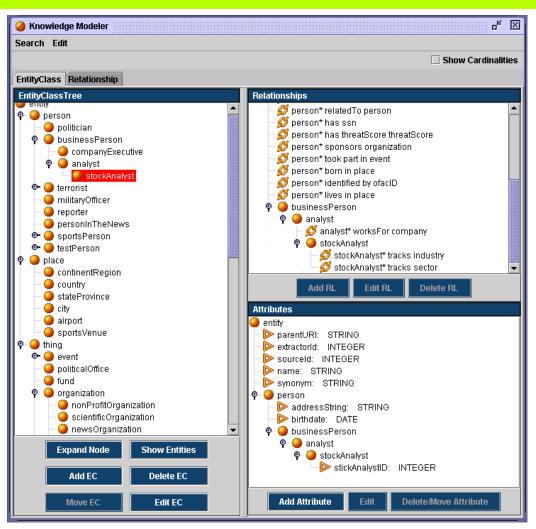


- This corresponds to the descriptional part (schema) of the Ontology
- Manually define Ontology structure (entity classes, relationship types, domain-specific and domain independent attributes)
- Configure parameters for attributes pertaining to indexing, lexical analysis, interface, etc.
- Existing industry-specific taxonomies like MESH (Medical), etc. can be reused or imported into the Ontology



Step 1: Ontology Model Creation

Create an Ontology Model using Semagix Freedom Toolkit GUIs (Cont.)



- This corresponds to the schema of the definitional part of the Ontology
- Manually define Ontology structure for knowledge (in terms of entities, entity attributes and relationships)
- Create entity class, organize them (e.g., in taxonomy)

```
e.g. Person

BusinessPerson

Analyst

StockAnalyst . . .
```

- Establish any number of meaningful (named) relationships between entity classes
 - e.g. Analyst <u>works for</u> Company
 StockAnalyst <u>tracks</u> Sector
 BusinessPerson <u>own shares in</u> Company...
- Set any number of attributes for entity classes e.g. Person

```
L Address <text>
L Birthdate <date>
```

StockAnalyst

StockAnalystID <integer>



Step 2: Knowledge Agent Creation

Create and configure Knowledge Agents to populate the Ontology

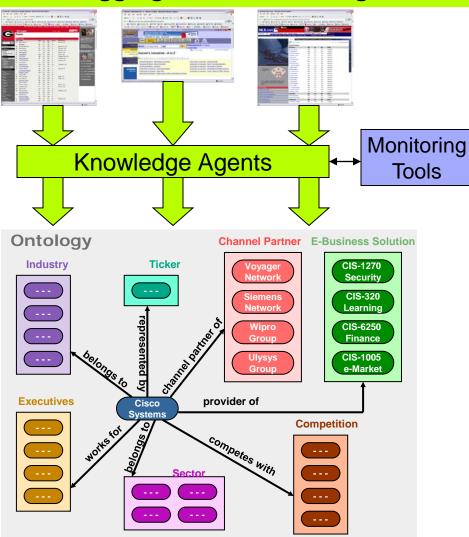


- Identify any number of trusted knowledge sources relevant to customer's domain from which to extract knowledge
 - Sources can be internal, external, secure/proprietary, public source, etc.
- Manually configure (one-time) the Knowledge Agent for a source by configuring
 - which relevant sections to crawl to
 - what knowledge to extract
 - what pre-defined intervals to extract knowledge at
- Knowledge Agent <u>automatically</u>) runs at the configured time-intervals and extracts entities and relationships from the source, to keep the Ontology up-to-date



Step 3: Automatic aggregation of knowledge

Automatic aggregation of knowledge from knowledge sources



- Automatic aggregation of knowledge at pre-defined intervals fo time
- Supplemented by easy-to-use monitoring tools
- Knowledge Agents extract and organize relevant knowledge into the Ontology, based on the Ontology Model
 - Tools for disambiguation and cleaning
- The Ontology is constantly growing and kept up-to-date

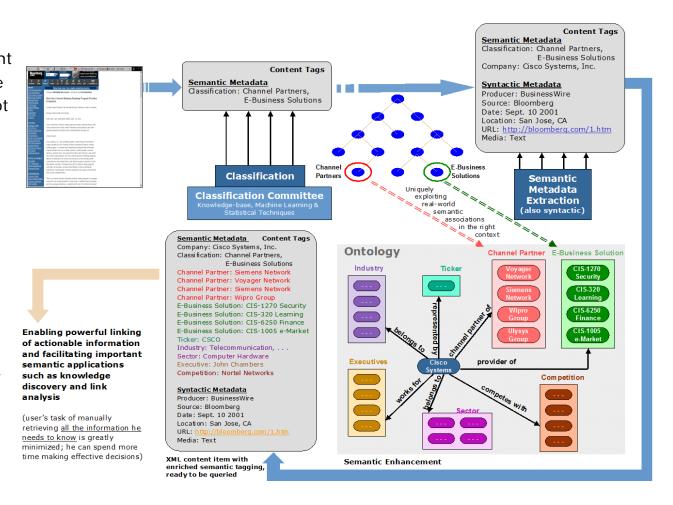
Semantic Enhancement Server

Semantic Enhancement

Server: Semantic Enhancement Server classifies content into the appropriate topic/category (if not already pre-classified), and subsequently performs entity extraction and content enhancement with semantic metadata from the Semagix Freedom Ontology

How does it work?

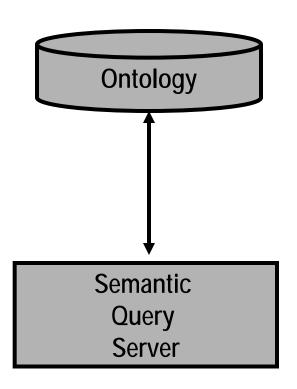
- Uses a hybrid of statistical, machine learning and knowledge-base techniques for classification
- Not only classifies, but also enhances semantic metadata with associated domain knowledge





Step 4: **Querying the Ontology**

Semantic Query Server can now query the Ontology



- Semantic Query Server can now perform in-memory complex querying on the Ontology and Metadata
 - Incremental indexing
 - Distributed indexing
 - High performance: 10M queries/hr; less than 10ms for typical search queries
 - 2 orders of magnitude faster than RDBMS for complex analytical queries
- Knowledge APIs provide a Java, JSP or an HTTP-based interface for querying the Ontology and Metadata

Ontology-based Semagix solutions

Equity Analysis Workbench

- Heterogeneous internal and extenral, push and pull content
- Automatic Classification, Semantic Information Correlation,
 Semantic (domain-specific search)

CIRAS - Anti Money Laundering:

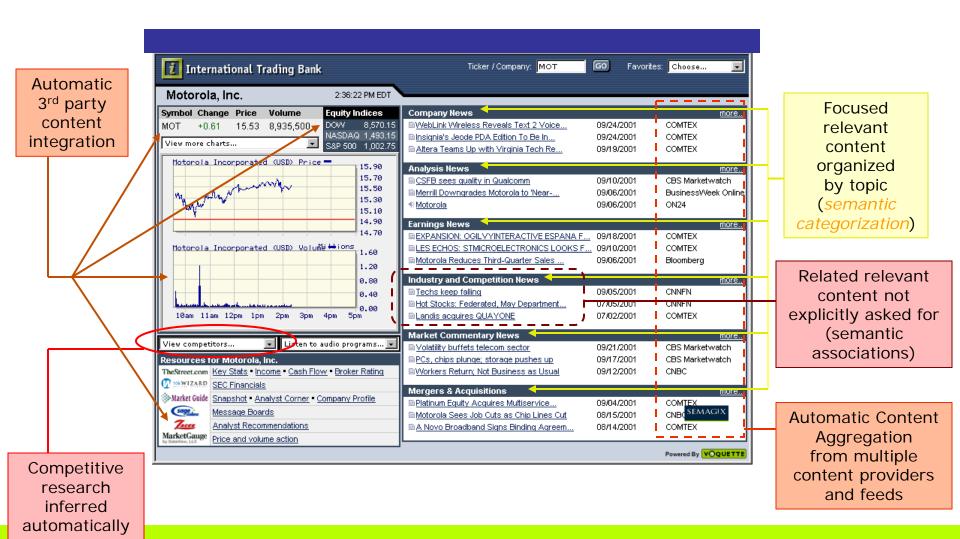
- Business issue: Optimisation of complex analysis from multiple sources
- Technology: Integration of process specific business insight from structured and unstructured information sources

APITAS – Passenger threat assessment

- Business issue: Rapid identification of high risk scenarios from vast amounts of information
- Technology: Managed high volume of information, speed of main memory indexed queries



Semantic Application Example – Analyst Workbench



CIRAS - Anti Money Laundering (Know Your Customer – KYC)



Fundamental Issues – Current Processes

Existing service bureau offerings created for different purpose — credit scoring

- Majority of content supplied not applicable to KYC unnecessary cost
- Rigid and static information require user interpretation elongation of process time
- Not specific enough to comply with new legislation non-compliance

Multiple manual checks against a variety of sources

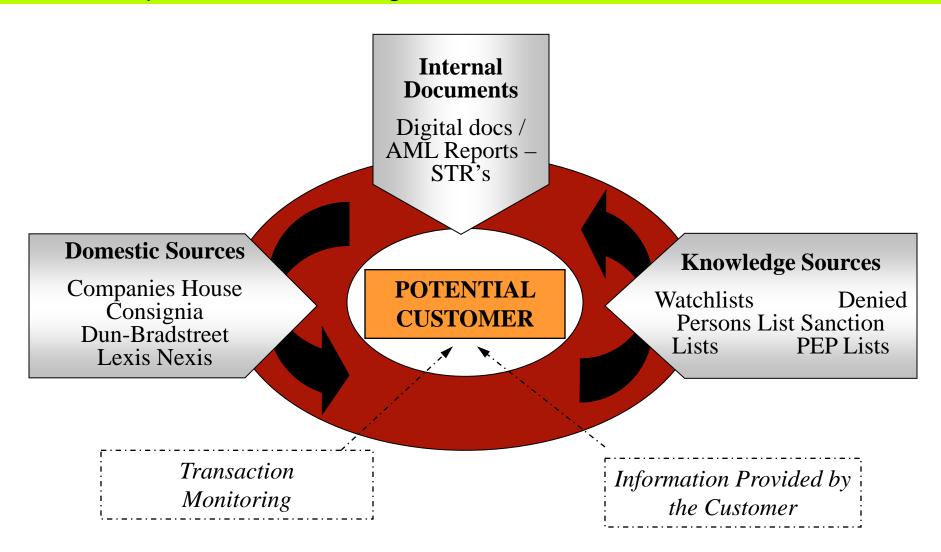
- Difficulty to link different pieces of information reduced effectiveness
- Checks are sequential and resource intensive Increase process time and cost
- Duplication of content increased subscription cost

Inability to implement domain-specific 'best practises'

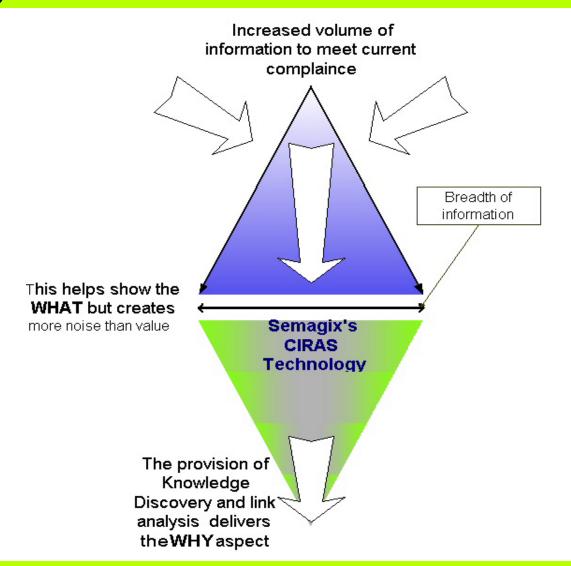
- Process knowledge resides with analysts variable quality of output
- Difficulty to fine-tune processes to specific domain inflexible process

Current processes are resource and time inefficient leading to inflexible and costly compliance

Constituent parts of 'reasonable grounds'



What vs. Why



What are the benefits

- Control compliance officers dictate the scale and scope of the checks made without incremental costs
- 2. Protects integrity of the company reputation and confidence are maintained through effective systems and controls
 - Comply with new legislations and regulations proceeds of crime act 2002 part 7, USA PATRIOT act

3. Cost

- Lower total cost for compliance with current and future legislation
- Lower content subscription and HR costs
- 4. Increased quality and efficiency of the compliance process
- 5. Integration into existing processes open standards enables the technology to be integrated into current KYC processes
- **6. Interoperability** provides integration across disparate legacy systems facilitating 'retrospective reviews' of customer bases







Semagix's Approach to KYC

This is achieved through:

- Risk weighting based on the underlying information and predefined criteria
 - Watchlist check
 - Link Analysis
 - ID Verification
- 2. Verification of the identity of a customer's name and address against domestic knowledge and content sources, includes:
 - What is already known about the customer
 - 3rd Party integration if required
 - Details of content relevant to 'knowing the customer'



Actionable Information

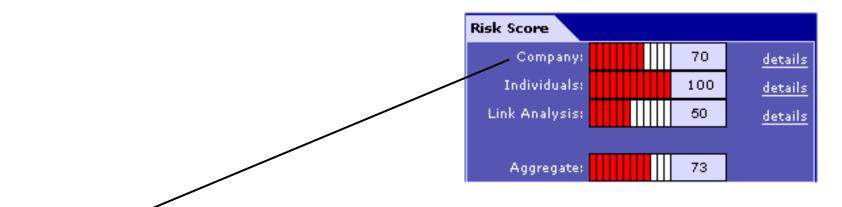


Aggregated risk represented by a customer

Summary of Capabilities

- Risk based approach to identification and verification
- Checks conducted against a wide variety of knowledge sources
- Integrates with existing processes
- Tailored for on-going and future requirements



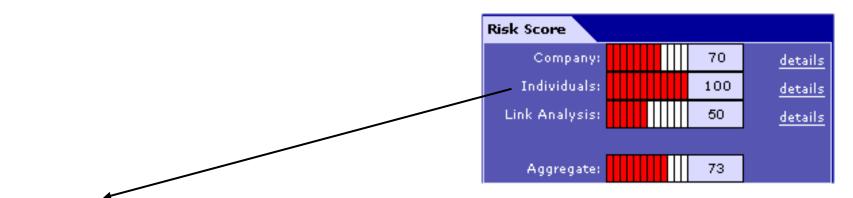


1. Company Analysis

Company Analysis - Details				
Score Component	Score	Reasons		
Watchlist/Sanction List Check	0.0			
Location Check	0.7	<u>Russia</u>		
Aggregate Score: 0.7				

- Cross references international and domestic watchlists
- Tailored to the operational environment
- Scheduled (every day) updates of the changes to lists

CIRAS's Components

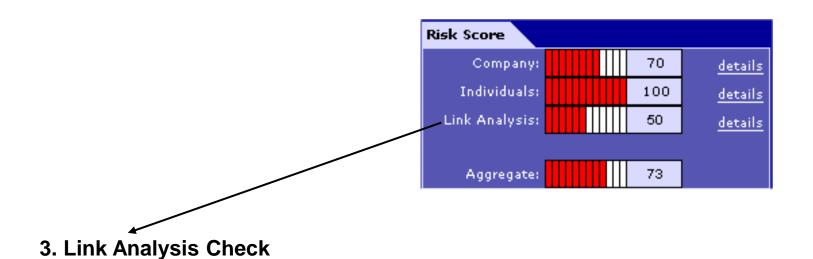


2. ID Verification

Analysis of Individuals - Details				
Score Component	Score	Reasons		
Watchlist/Sanction List Check	1.0	Richard Pott		
Company/Organisation	0.0			
Aggregate Score: 1.0				

- Provides an indication as to the risk posed by individuals associated with the company
- Allows navigation into possible causes of 'false positive's

CIRAS's Components

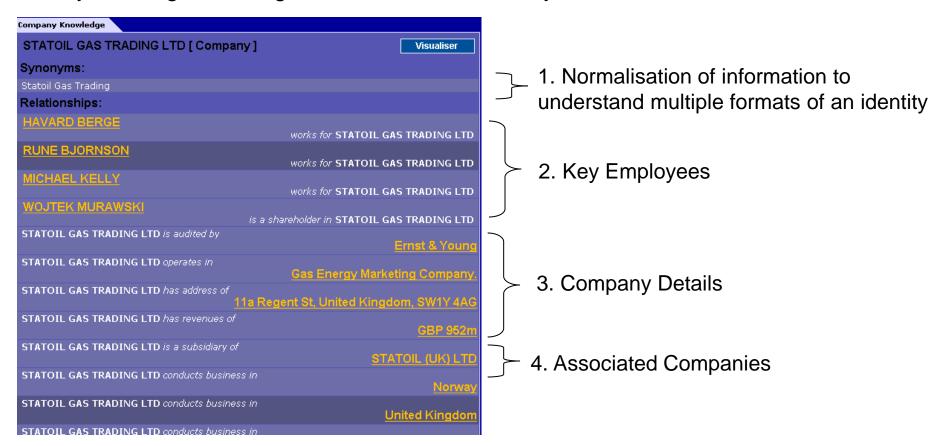


Link Analysis - Details			
Score Component	Score	Reasons	
Metabase Check	0.0		
Organisation Check	1.0	Akida Bank	
Aggregate Score: 0.0			

- Identification and verification of relationships customer holds with other entities (organisations, people etc)
- Flags high-risk transaction flows
- References internal reports held



Provision of 'knowledge' already held about a prospect and provides the ability to navigate through each 'instance' to verify information





External content, from multiple sources, in any format relevant to 'knowing the customer'

Internal content, previous KYC checks undertaken, STR reports filed and transaction monitoring alerts relevant to the customer in question





Current applications of the technology

- CIRAS Anti Money Laundering
- Passenger Threat Assessment System

External demo page



About Semagix

Semagix, through a patented semantic approach to Enterprise Information Integration (EII), allows enterprises to integrate and extract insights from their structured and unstructured information assets in order to conceive and develop smarter business processes and applications

